

CSC 212

Data Structures and Abstractions (Spring 2025) C++ Review, Memory, and Pointers

Prof. Marco Alvarez, University of Rhode Island

Context

machine code

```
10110100 10110111 00101011
00011010 00010100 10111011
10001000 11110111 00101000
10101010 00101101 00010001
01010010 11101100 11010001
10010100 10010100 00100000
00000000 10100001 00110001
10101001 00010101 00101010
00100100 10010100 01110001
11110101 11101011 00101111
01010010 10000101 11111110
10101010 00101101 00010001
01010010 11101100 11010001
10010100 10010100 00100000
00000000 10100001 00110001
10101001 00010101 00101010
00100100 10010100 01110001
11110101 11101011 00101111
01010010 10000101 11111110
00101001 00000000 00000000
00000000 00000000 01010000
00001010 00010101 00101010
00101010 00100100 10011111
```

assembly

```
.equ STDOUT, 1
.equ SVC_WRITE, 64
.equ SVC_EXIT, 93

.text
.global _start
_start:
    stp x29, x30, [sp, -16]!
    mov x0, #STDOUT
    ldr x1, =msg
    mov x2, 13
    mov x8, #SVC_WRITE
    mov x29, sp
    svc #0
    ldp x29, x30, [sp], 16
    mov x0, #0
    mov x8, #SVC_EXIT
    svc #0

msg: .ascii "Hello World!\n"
    .align 4
```

C++

```
#include <iostream>

int main () {
    std::cout << "Hello World!"
    << std::endl;
}
```

Python

```
print('Hello World')
```

→ increasing abstraction →

https://www.uvm.edu/~cbcafier/cs1210/book/02_programming_and_the_python_shell/programming.html

To illustrate the potential gains from performance engineering, consider multiplying two 4096-by-4096 matrices. Here is the four-line kernel of Python code for matrix-multiplication:

```
for i in xrange(4096):
    for j in xrange(4096):
        for k in xrange(4096):
            C[i][j] += A[i][k] * B[k][j]
```

Version	Implementation	Running time (s)	GFLOPS	Absolute speedup	Relative speedup	Fraction of peak (%)
1	Python	25,552.48	0.005	1	—	0.00
2	Java	2,372.68	0.058	11	10.8	0.01
3	C	542.67	0.253	47	4.4	0.03
4	Parallel loops	69.80	1.969	366	7.8	0.24
5	Parallel divide and conquer	3.80	36.180	6,727	18.4	4.33
6	plus vectorization	1.10	124.914	23,224	3.5	14.96
7	plus AVX intrinsics	0.41	337.812	62,806	2.7	40.45

From: "There's plenty of room at the Top: What will drive computer performance after Moore's law?"

Program execution approaches

• Compilation

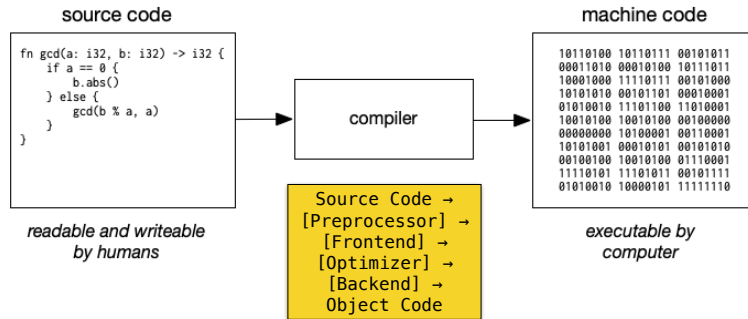
- high level source **translated** into another language
 - most of the time into a low-level language
- as code is translated at once, compilers can perform **optimizations** to make the code more efficient, resulting in faster execution (higher performance)
- e.g. C/C++ compilers

• Interpretation

- 'executing' a program directly from source
 - read code line by line, translate it into machine code, and execute
 - any language can be interpreted
- preferred when performance is not critical
- e.g. Javascript


Compiling programs (simplified)

- Typically, “compiling” a program refers to the process of generating machine code from source code
- the process takes several steps: **compile**, **assemble**, **link**



https://www.uvm.edu/~cbcafier/cs1210/book/02_programming_and_the_python_shell/programming.html

Compiling/linking/running C programs



From Computer Organization and Computer Design: The Hardware/Software Interface

What is the output?

```
#include <iostream>  
  
int main() {  
  int d = 42;  
  int o = 052;  
  int x = 0x2a;  
  int X = 0X2A;  
  int b = 0b101010; // C++14  
  
  std::cout << d << " " << o << " " <<  
    x << " " << X << " " << b << std::endl;  
  
  return 0;  
}
```

Range of values (fundamental types)

Data type	Size	Format	Value range
character	8	signed	-128 to 127
		unsigned	0 to 255
integer	16	signed	-32768 to 32767
		unsigned	0 to 65535
	32	signed	-2,147,483,648 to 2,147,483,647
		unsigned	0 to 4,294,967,295
	64	signed	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
		unsigned	0 to 18,446,744,073,709,551,615

<https://en.cppreference.com/w/cpp/language/types>

Integral types

Type specifier	Equivalent type	Width in bits by data model				
		C++ standard	LP32	ILP32	LLP64	LP64
signed char	signed char	at least 8	8	8	8	8
unsigned char	unsigned char					
short	short int	at least 16	16	16	16	16
short int						
signed short						
signed short int						
signed short int	unsigned short int					
unsigned short						
unsigned short int	int	at least 16	16	32	32	32
int						
signed						
signed int						
unsigned	unsigned int					
unsigned int						
long	long int	at least 32	32	32	32	64
long int						
signed long						
signed long int						
unsigned long	unsigned long int					
unsigned long int						
long long	long long int (C++11)	at least 64	64	64	64	64
long long int						
signed long long						
signed long long int						
unsigned long long	unsigned long long int (C++11)					
unsigned long long int						

Memory organization

Memory organization

- Memory as a **byte array**
 - used to store **data and instructions** for computer programs
 - contiguous sequence of bytes
 - each byte individually accessed via a **unique address**
- Memory address
 - **unique** numerical identifier for each byte in memory
 - **pointer** variables store memory addresses
 - provides indirect access to data stored at that location

Memory organization


- Data representation in memory
 - variables stored as byte sequences
 - interpretation depends on type
 - integers, floating-point numbers, characters, etc.
- OS provides private address space to each **“process”**
 - process: a program being executed
 - address space: enormous arrays of bytes visible to the process
 - typically implemented through virtual memory

Pointers

Variables and pointers

- Every variable exists at a **memory address** (regardless of **scope**)
 - memory address corresponds to a unique location
- The compiler translates names to addresses when generating machine code

A **pointer** is a variable that stores the address of another variable



Pointers

- Must be declared before use
 - **pointer type** must be specified
- Pointer operators
 - **address-of operator**: get memory address of variable/object

&

- **dereference operator**: get value at given memory address

*

Null pointers and arrays

- The NULL value (`0x00000000`)
 - represents the absence of value
 - reading/writing with a null value can generate a **segmentation fault** signal
 - in C++, it is safer to use `nullptr` (keyword)
- Pointers and arrays
 - arrays decay to pointers (to the first element) in most contexts, but they are not themselves pointers
 - array names provide the address of the first element, can't be treated as variables

Declaring pointers

```
// can declare a single
// pointer (preferred)
int *p;

// can declare multiple
// pointers of the same type
int *p1, *p2;

// can declare pointers
// and other variables too
double *p3, var, *p4;
```

Pointer operators

```
int main() {
    int var = 10;
    int *ptr;
    ptr = &var;
    *ptr = 20;

    // ...

    return 0;
}
```

32-bit words

Address	Value	Variable
...		
0x91340A08		
0x91340A0C		
0x91340A10		
0x91340A14		
0x91340A18		
0x91340A1C		
0x91340A20		
0x91340A24		
0x91340A28		
0x91340A2C		
0x91340A30		
0x91340A34		
...		

Pointer operators

```
int main() {
    int temp = 10;
    int value = 100;
    int *p1, *p2;

    p1 = &temp;
    *p1 += 10;

    p2 = &value;
    *p2 += 5;

    p2 = p1;
    *p2 += 5;

    return 0;
}
```

32-bit words

Address	Value	Variable
...		
0x91340A08		
0x91340A0C		
0x91340A10		
0x91340A14		
0x91340A18		
0x91340A1C		
0x91340A20		
0x91340A24		
0x91340A28		
0x91340A2C		
0x91340A30		
0x91340A34		
...		

Pointers and functions

```
void increment(int *ptr) {
    (*ptr) ++;
}

int main() {
    int var = 10;

    increment(&var);
    increment(&var);

    // ...

    return 0;
}
```

32-bit words

Address	Value	Variable
...		
0x91340A08		
0x91340A0C		
0x91340A10		
0x91340A14		
0x91340A18		
0x91340A1C		
0x91340A20		
0x91340A24		
0x91340A28		
0x91340A2C		
0x91340A30		
0x91340A34		
...		

Pointer arithmetic

- Can add values to pointers
 - treats addresses as unsigned integers
- Must be careful !
 - `p+1` adds the size of pointed variable
 - `p+1` does NOT add 1 “byte”
- Can use pointer arithmetic for array traversal

`a[i]` is equivalent to `*(a+i)`

Changing a pointer inside a function

```
#include <stdio.h>

void seek(int *p, int key, int n) {
    for (int i = 0 ; i < n; i++) {
        if (*p == key) {
            return;
        }
        p ++;
    }
}

int main() {
    int data[] = {1, 2, 3, 4, 5};
    int *p = data;

    seek(data, 3, 5);
    printf("%d\n", *p);

    return 0;
}
```

does it work?

Using double pointers

```
// function to search for a key in an array
// - pointer to an array of integers
// - an integer key
// - an integer n, the number of elements
```

```
void seek(int **p, int key, int n) {
    for (int i = 0 ; i < n; i++) {
        if (**p == key) {
            return;
        }
        (*p) ++;
    }
}
```

Using double pointers

```
int main() {
    int data[] = {1, 2, 3, 4, 5};
    int *p = data;

    seek(&p, 3, 5);
    printf("%d\n", *p);

    return 0;
}
```

Python Tutor: Visualize code in Python, JavaScript, C, C++, and Java

C (C17 + GNU extensions)

[known limitations](#)

```
6 // - an integer key
7 // - an integer n, the number of elements in the
8 void seek(int **p, int key, int n) {
9     for (int i = 0 ; i < n; i++) {
10         if (**p == key) {
11             return;
12         }
13     }
14     (*p)++;
15 }
16
17 int main() {
18     int data[] = {1, 2, 3, 4, 5};
19     int *p = data;
20
21     seek(&p, 3, 5);
22     printf("%d\n", *p);
23
24     return 0;
25 }
```

[Edit this code](#)

⇒ line that just executed

➔ next line to execute

<< First

< Prev

Next >

Last >>

Step 9 of 17

Print output (drag lower right corner to resize)

Stack Heap

main

data

0	1	2	3	4
int	int	int	int	int
1	2	3	4	5

p

pointer to int

seek

p

pointer to int*

key

int

3

n

int

5

i

int

0

C/C++ details: none [default view]