

CSC 212: Data Structures and Abstractions

Balanced trees (part 2)

Prof. Marco Alvarez

Department of Computer Science and Statistics
University of Rhode Island

Spring 2025



Red-black trees

Insertion

Steps

- ✓ insert the new node following standard BST rules
- ✓ color the new node **red** (required to maintain the *root-to-null* rule)
- ✓ if parent is **black**, terminate (forms 3-node or 4-node)
- ✓ if parent is **red**, resolve the *red-red* violation by propagating fixes upward
- ✓ finally, make sure the root is **black**

Violation resolution

- ✓ apply **recoloring** (preserves structure) and/or **rotations** (preserves order)
 - rotations are used when recoloring alone cannot restore properties

3

Rotations


- A **rotation** is a $O(1)$ -time local operation that preserves the BST order property while modifying structure

Right rotation at node y

- ✓ requires y 's left child x to be *non-null*
- ✓ elevates x to become the subtree root
- ✓ y becomes x 's right child
- ✓ x 's original right child becomes y 's left child

Left rotation at node x

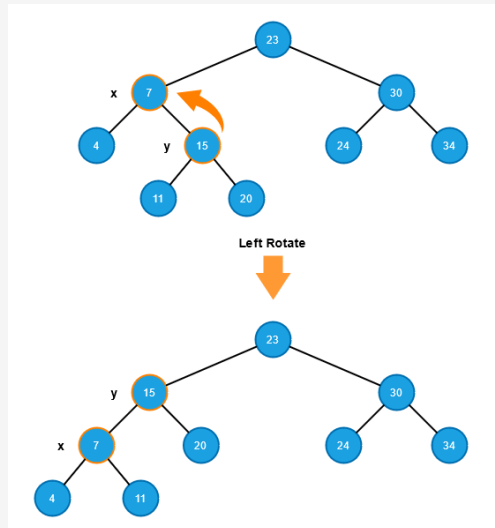
- ✓ requires x 's right child y to be *non-null*
- ✓ elevates y to become the subtree root
- ✓ x becomes y 's left child
- ✓ y 's original left child becomes x 's right child



4

Example: left rotation

left-rotate(x)




<https://www.formosa1544.com/2021/04/30/build-the-forest-in-python-series-red-black-tree/>

5

Example: right rotation

right-rotate(x)



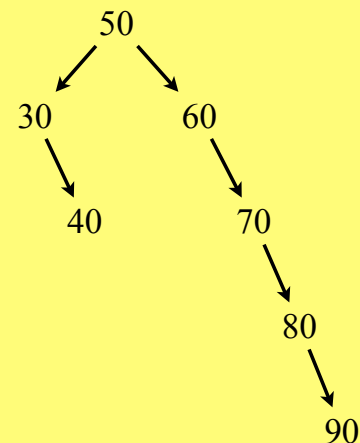
<https://www.formosa1544.com/2021/04/30/build-the-forest-in-python-series-red-black-tree/>

6

Practice

• Perform the following operations in sequence

- ✓ rotate-left(70)
- ✓ rotate-left(50)
- ✓ rotate-left(30)
- ✓ rotate-right(50)



7

Insertion (detailed steps)

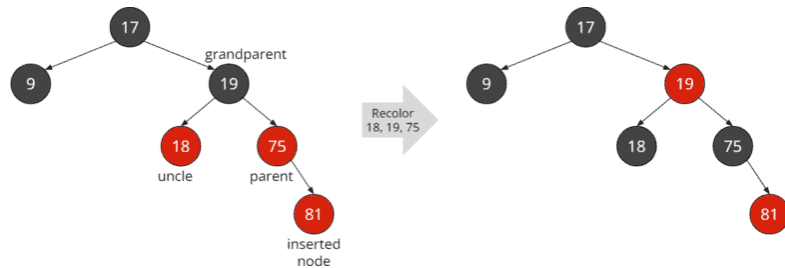
- Insert the new node following standard BST rules
 - ✓ color the new node **red**
 - ✓ apply the appropriate case resolution

• Case resolution

- ✓ **Case 1:** parent of new node is **black**
 - no *red-red* violation occurs
 - the tree remains valid without modification
- ✓ **Case 2:** parent and uncle of new node are **red**
 - recolor parent and uncle to black
 - recolor grandparent to red
 - recursively validate the grandparent as if newly inserted

8

Example: case 2



consider 3 other analogous (sub)cases

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

9

Insertion (detailed steps)

• Case resolution (cont.)

✓ Case 3: triangle formation

- condition: parent is **red**, uncle is **black** (or null)
- action: apply rotation at parent to convert to line formation
 - if new node is right child of left parent, perform left rotation at parent
 - if new node is left child of right parent, perform right rotation at parent
- proceed to case 4

✓ Case 4: line formation

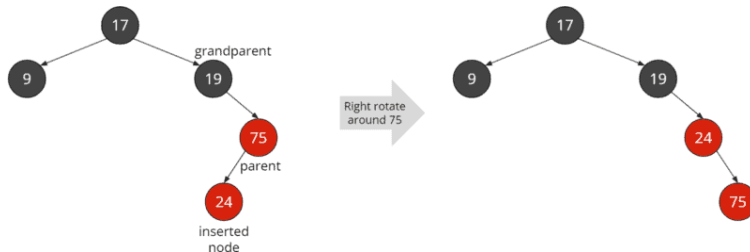
- condition: parent is **red**, uncle is **black** (or null)
- action: rotate at grandparent and then swap colors of original parent and grandparent
 - right rotation if line extends left
 - left rotation if line extends right
- this case terminates the rebalancing process

• Final step

- ✓ after all case resolutions, ensure the root node is colored **black**

10

Example: case 3 (triangle)

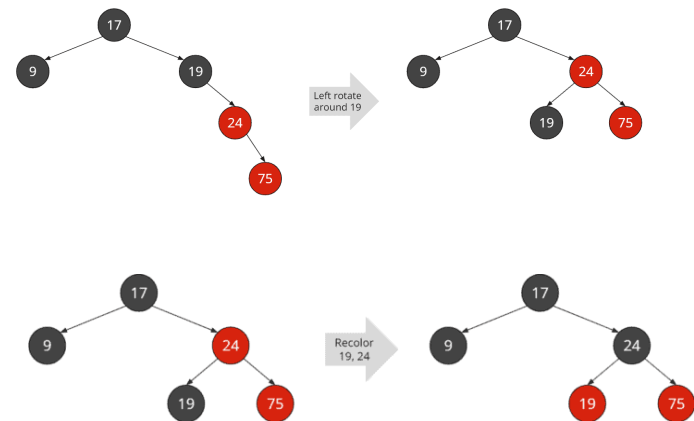


consider 1 other analogous (sub)case

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

11

Example: case 4 (line)



consider 1 other analogous (sub)case

<https://www.happycoders.eu/algorithms/red-black-tree-java/>

12

Practice

- Insert the following keys into a RB-tree
 - 10, 20, 30, 40, 50, 15, 25, 35, 45

13

Final remarks

- Theoretical Equivalence**
 - due to the isometry between red-black trees and 2-3-4 trees, the maximum height of a red-black tree with n nodes: $2 \log_2(n + 1)$
- Other operations**
 - search**: identical to binary search trees
 - delete**: similar to insert (not covered in lecture, standard libraries provide optimized implementations)
- C++ Implementation**
 - STL containers `std::set` and `std::map` typically use red-black trees

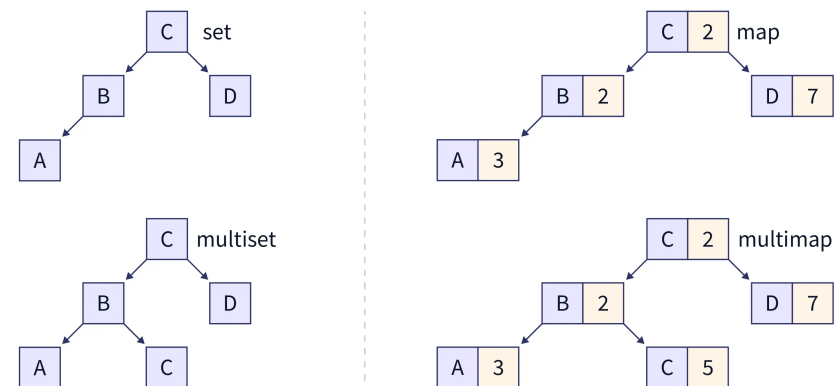
14

Analysis

Data Structure	Worst-case			Average-case			Ordered?
	insert at	delete	search	insert at	delete	search	
sequential (unordered)	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	No
sequential (ordered) binary search	$O(n)$	$O(n)$	$O(\log n)$	$O(n)$	$O(n)$	$O(\log n)$	Yes
BST	$O(n)$	$O(n)$	$O(n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
2-3-4	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes
Red-Black	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	$O(\log n)$	Yes

15

Ordered associative containers (STL)



<https://www.scaler.com/topics/cpp/containers-in-cpp/>

16

std::set

Defined in header `<set>`

```
template<
    class Key,
    class Compare = std::less<Key>,
    class Allocator = std::allocator<Key>
> class set; (1)

namespace pmr {
    template<
        class Key,
        class Compare = std::less<Key>
    > using set = std::set<Key, Compare, std::pmr::polymorphic_allocator<Key>>; (2) (since C++17)
}
```

`std::set` is an associative container that contains a **sorted set of unique objects of type `Key`**. Sorting is done using the key comparison function `Compare`. Search, removal, and insertion operations have logarithmic complexity. Sets are usually implemented as [Red-black trees](#).

Everywhere the standard library uses the `Compare` requirements, uniqueness is determined by using the equivalence relation. In imprecise terms, two objects `a` and `b` are considered equivalent if neither compares less than the other: `!comp(a, b) && !comp(b, a)`.

`std::set` meets the requirements of [Container](#), [AllocatorAwareContainer](#), [AssociativeContainer](#) and [ReversibleContainer](#).

```
#include <iostream>
#include <set>

int main() {
    std::set<int> example{1, 2, 3, 4};

    for (int x : {2, 5}) {
        if (example.contains(x))
            std::cout << x << ": Found\n";
        else
            std::cout << x << ": Not found\n";
    }
}
```

17

std::map

Defined in header `<map>`

```
template<
    class Key,
    class T,
    class Compare = std::less<Key>,
    class Allocator = std::allocator<std::pair<const Key, T>>
> class map; (1)

namespace pmr {
    template<
        class Key,
        class T,
        class Compare = std::less<Key>
    > using map = std::map<Key, T, Compare,
        std::pmr::polymorphic_allocator<std::pair<const Key, T>>>; (2) (since C++17)
}
```

`std::map` is a sorted associative container that contains key-value pairs with unique keys. Keys are sorted by using the comparison function `Compare`. Search, removal, and insertion operations have logarithmic complexity. Maps are usually implemented as [Red-black trees](#).

Iterators of `std::map` iterate in ascending order of keys, where ascending is defined by the comparison that was used for construction. That is, given

- `m`, a `std::map`
- `it_l` and `it_r`, dereferenceable iterators to `m`, with `it_l < it_r`.

`m.value_comp()(*it_l, *it_r) == true` (least to greatest if using the default comparison).

Everywhere the standard library uses the `Compare` requirements, uniqueness is determined by using the equivalence relation. In imprecise terms, two objects `a` and `b` are considered equivalent (not unique) if neither compares less than the other: `!comp(a, b) && !comp(b, a)`.

`std::map` meets the requirements of [Container](#), [AllocatorAwareContainer](#), [AssociativeContainer](#) and [ReversibleContainer](#).

18